

The Invisible Architect: (Mis)Perceptions of Agency Around GenAI Use in Recruiting Workflows

Sajel Surati
New York University
New York, New York, USA
sajel.s@nyu.edu

Rosanna Bellini
New York University
New York, New York, USA
bellini@nyu.edu

Emily Black
New York University
New York, New York, USA
emilyblack@nyu.edu

Abstract

When generative AI (genAI) systems are used in high-stakes decision-making, its recommended role is to aid, rather than replace, human decision making. However, there is little empirical exploration of how professionals making high-stakes decisions, such as those related to employment, perceive their agency and level of control when working with genAI systems. Through interviews with 22 recruiting professionals, we investigate how genAI perceptibly and subtly influences control over everyday workflow and even individual hiring decisions. Our findings highlight a pressing conflict; while recruiters believe they have final authority across the recruiting pipeline, genAI has become an invisible architect that shapes the foundational building blocks of information used for evaluation, from defining a job to determining good interview performances. We conclude by discussing the implications for such findings for responsible and perceptible genAI use in hiring contexts. We pose that such reliance on genAI tools can lead to a loss of agency, inhibiting proper bidirectional human-AI alignment.

CCS Concepts

• **Social and professional topics** → *Computing / technology policy*.

Keywords

Algorithmic Decision-making, Generative AI, Recruitment, Human-Computer Interaction, Behavioral Agency

1 Introduction

Generative AI systems (genAI) are transforming the workplace, with studies reporting adoption rates up to 80% [2, 44]. Recruitment is an exemplar of this shift: major platforms such as LinkedIn, alongside numerous startups, now integrate genAI across the recruitment funnel; from sourcing and screening to automating the interviewing process itself.

In high-stakes decision-making settings like hiring, scholars and policymakers emphasize the need for “AI-assisted” decision-making; asserting that AI should support, not replace human judgment [22, 38]. But how realistic is maintenance of human agency in practice? While prior work has both problematized and praised “human in the loop” or “AI assisted” decision-making through controlled experiments [20, 26–28, 55], there is little empirical work investigating how practitioners actually navigate incorporating AI—especially generative AI—into their decision-making on the ground. We use the field of recruitment and hiring as a case study

to examine how agency is experienced and redistributed in AI-mediated hiring processes — both on the level of making individual hiring decisions, and systemically, regarding the adoption of or resistance to using genAI in their workflows. Specifically, we investigate the following research question: *To what extent do recruiters feel they maintain agency in their decision-making?*

We investigate this question through interviews with 22 recruiters across six industry sectors, identifying a widening gap between perceived professional autonomy and the reality of AI-driven workflows. The mechanisms and implications of this phenomenon form our main contributions.

On an individual level (Section 4.1), we find that while recruiters claim ultimate authority—often asserting that AI does not make decisions or is relegated to low-level tasks—this presents an incomplete picture. Instead, our findings characterize genAI as an *invisible architect*. By defining job qualifications, creating interview rubrics, and summarizing candidate interactions, genAI structures the decision space and exerts a profound, unacknowledged influence on hiring — often before humans realize a decision is being made.

Beyond this contribution, to map how genAI systems enter hiring workflows, we perform a functionality analysis of genAI hiring tools mentioned in interviews (Appendix B), and discuss how our findings inform paths towards reclaiming human agency in recruiting workflows through AI governance and future research in Section 5. We conclude with implications for bidirectional human-AI alignment, suggesting that further work engage with disruptive design techniques to return agency back to decision-makers.

2 Background and Related Work

Recruiting refers to the process of identifying, attracting, and selecting individuals to fill specific job roles within a team or a company. This may be performed by recruiters who are either in-house (often located in human resource departments), contracted out to an external agency, or performed by a current employee (a *hiring manager*). A recruitment pipeline consists of five chronological stages: job design, sourcing, screening, interviewing, and offering. First, a job is designed or ‘themed’ whereby a person (typically a hiring manager) defines the responsibilities for a position. Recruiters can either advertise the position to attract candidates (inbound), or approach viable candidates (outbound) with a proposition to apply. Applicants are then screened, where their application materials are evaluated against the position’s requirements and organizational factors; those who pass may be invited to an interview to evaluate their qualifications and ‘fit’ across different teams or increasing

staff seniority level. A hiring manager may then select a candidate to extend a job offer.

Algorithmic Hiring Decisions. Our work studies the impacts of AI and genAI tools on individual’s perception of their agency in recruitment workflows. Prior literature highlights potential ethical concerns regarding AI in hiring, largely centered around the potential for discrimination. Several papers provide technical audits have quantified rates of discrimination in both traditional AI [9, 57] and genAI [3, 4, 6, 24, 47, 58], while socio-technical audits have scrutinized claims made by companies developing recruitment tools [49]. Closely related to our work, Min et al. [39] interviewed recruiting professionals regarding traditional in screening processes, though their focus remains on barriers to adoption. Similarly, Wilson et al. [59] perform a user study on how humans are influenced by AI-generated candidate interview recommendations, and demonstrate that humans are indeed influenced by AI-based recommendations. However, this study was not interview-based did not include human resource professionals, and focused on isolated recommendations rather than the broader recruitment pipeline.

Behavioral Agency in Human-Machine Interactions. Our work builds upon regarding human agency in AI-assisted decision-making teams. In this work, we focus on *behavioral agency* (derived from *agere* “to set in motion, drive forward; to do, perform” [1]), which focuses on an individual’s or group’s capacity to shape decision-making. Behavioral agency been explored across a wealth of complex sociotechnical systems [16, 49, 53, 65], being so central to well-designed systems that developers have long been encouraged to “support internal locus of control” [52], despite the tension between human and machine agency [53]. Human reliance on AI is driven by internal factors, such as trust and affinity for technology [16, 37], alongside system traits, like perceived fairness or task simplicity [54]. Thus, investigations into human-AI interactions suggest that decision-making is characterized by relational agency, where a human’s capacity to act is constrained by an AI’s architecture [42, 45].

GenAI complicates behavioral agency through introducing problems into the heart of decision-making. For instance, LLMs can display “social sycophancy”; human-like reasoning and mimic authoritative tones, whereby a model prioritizes saving face over providing critical or truthful feedback [17, 43]. Bućinca [14] and Lai [38] thus caution that the true locus of decision-making authority often resides in pre-computational stages; how an AI frames, filters, and prioritizes data before a practitioner intervenes. When users rely on AI systems, problems involving governance [8], high-stakes errors [63], and deskilling [18, 36, 64] become ever more present [31]. By interviewing recruiters, we hope to gain insight into users’ perceived reliance on genAI in high-stakes decision-making.

3 Methodological Approach

Our study adopts a mixed methods approach, whereby we conducted a series of semi-structured interviews with 22 recruiters across the space of four months (Sept – Dec 2025) on their acceptance and use of such tools in their work practices (Section 4).

Semi-structured interviews. Given the rapidly evolving nature of genAI integration with recruitment [41], we chose to use semi-structured interviews for both epistemological and practical reasons.

We hoped that the conversational atmosphere of a semi-structured interview could help build rapport with participants, leading to candid disclosures.

Recruitment Strategy. Participants were recruited using a multi-pronged sampling strategy. First, we conducted purposive sampling via LinkedIn ($N=19$), targeting employees at organizations via the active job postings search feature to increase the likelihood of reaching potential interviewees who were engaged in hiring. To supplement this approach, we also used convenience sampling ($N=1$) by contacting professional acquaintances who had previously described using genAI in their recruitment workflows. Finally, the study also attempted to use referral-based (i.e., snowball) sampling ($N=2$) by encouraging participants to disseminate our study details within their professional networks.

Interview Process. Our interview protocol was designed around four topic areas: introductory conversation, grounding the context, identifying how use of genAI has been used in hiring, their perception of agency over choices regarding their work, and the future potential of AI. Interested readers can find a copy of our interview protocol in Appendix C.

Data were collected via synchronous, semi-structured interviews conducted over Zoom video-conferencing software. Data that were collected were automatically transcribed using Zoom auto-transcribe before being manually verified and anonymized.

Analysis. Our interview data were analyzed using a reflexive thematic analysis approach [11] using the qualitative analysis software *Atlas.ti* [25]. Following familiarization by conducting a close reading of the transcripts, we used a hybrid approach of inductive and deductive coding. We drafted a initial set of deductive codes ($N=14$) in accordance with our research questions (Section 1), which were then supplemented with a round of inductive open coding from two transcripts, resulting in 67 additional codes, for a total of 81 initial codes. The codebook was refined further through two subsequent cycles with synchronous meetings for coordinating disagreements for a final codebook of 54 codes. Inter-rater reliability of Krippendorff’s Alpha were calculated as a mechanism for consensus-building, at 0.80 agreement (“Very Good” [35]). See Figure 3 in Appendix (D) for a copy of our final codebook.

Ethics and Limitations. This study reviewed and judged to be exempt by the New York University Institutional Review Board (IRB). Participants were sent an information sheet and a consent form prior to the interview, and written informed consent were taken prior to the interview with participants being asked to re-confirm their consent verbally at the start of the session. Audio recording was voluntary and de-identified note taking was offered as an alternative. Participants were provided a \$50 USD honorarium as remuneration for their time and contributions.

Given the potential for social desirability bias [11], whereby participants could feel pressure to report policy-compliant or ‘ethical’ use of genAI, we designed our protocol to utilize a normalization framing to reduce feelings of discomfort. To further protect our participants from potential retaliation by employers, we chose to not report company names or sizes, report agnostic job titles, and report participant demographics in aggregate where possible (see Appendix 1). We also report our findings as aggregated themes

whereby idiosyncratic details (e.g., software, team structure) are withheld [10].

Our study also has the following limitations. As our participants report their genAI use retrospectively, this may have limited their ability to accurately recall specific actions. We also encountered a minority of participants ($N=3$) who omitted specific details of their work practice citing company policies. Finally, only interviewing participants in the US means that the international perspective is underexplored, which we look forward to in future work.

4 Findings

We present our findings: how genAI subtly erodes human agency through acting as an *invisible architect* in recruitment decision-making and that despite this erosion, recruiters still feel in control over the hiring process.

4.1 RQ: To what extent do recruiters feel they maintain agency in their decision-making?

Recruiters overwhelmingly reported feeling that they alone were in control of the hiring decisions they made. However, our interviews showed that genAI systems could potentially have more influence over these choices than many participants gave them credit for. We identified that genAI systems were predominantly used to *frame* decisions: our interviews shared many descriptions of genAI being used to both create assessment schemes of candidate quality, from job descriptions, interview questions, and rubrics, as well as process the information that went into presenting the candidates in the first place; such as summarizing resumes, interview performance and more. Thus, our analysis identified that genAI acted as an *invisible architect*; an authority on organizing the framework in which people make decisions, and the information used in decisions, without the necessary visibility that could make participants aware of how their behavior was being shaped. As P_3 shared,

"[AI] helps us to really identify the people we've been talking to, and boil down to the different patterns that exist... And then using that, it creates, like, an ideal profile, and then from there, you can create a job description. And then it also says, hey, based on this, here is an ideal interview plan. And so it kind of helps you develop the entire operations that way, too, but also make decisions, because you are going to lean on it for some ... Some data, some information, so on and so forth." (P_3)

Here, P_3 acknowledges that by developing most of the materials used to assess candidates and processing all the information used to go into these assessments, genAI does actually influence the decision-making process itself because you must "*lean on it for... information*."

4.1.1 Recruiters feel in control and that AI only helps with low-level tasks. More broadly, we observed this dichotomy of perceived control and subtle AI influence through two co-occurring narratives. On the one hand, participants commonly reported feeling in control of hiring decisions, and chalked up the work of AI assistance, especially generative AI assistance, to low-level tasks. For example, regarding agency, P_{10} stated that "*I'm the deciding factor, right?*",

and P_9 stated that "*Bottom line, we don't use AI to make any, like, real decisions in hiring.*" On the other hand, as we demonstrate in Section 4.1.2, a considerable amount of decision-making power was passed on to AI systems via inherent trust of and deference to AI outputs, and using AI-generated framing and AI-summarized information to make decisions.

Part of this narrative that AI was not making decisions seemed to stem from the idea that AI was only doing low-level work: for example, P_{18} stated that the internal generative AI tool in her office that she used to do the initial screening on candidates was "*It's basically like a filter. You know, it's filtering it so that it doesn't waste my time.*" P_{21} stated that the percentage match between a candidates' resume and job description that she prompts ChatGPT to come up with is "*A helpful tip*". However, some of the descriptions of AI use had internal contradictions about its role and power: for example, P_5 said their AI interview tool "*Doesn't make decisions... It's completely just transcript only,*" despite noting that they used it to organize and highlight information from the interview.

Part of internal contradiction about the amount of control given over to AI may be due to the fact that many participants reported that AI should not be making final decisions about a candidate. As P_9 shared, "*I don't want to make any decisions with AI.*" However, despite this mentality, in many ways throughout the recruitment process, control was given to AI systems, in that their output was rarely doubted or double checked. As P_{18} said about her AI filtering system, that she "*[doesn't] have to necessarily double-check*" because she can trust the systems' output. This feeling of control even extended to recruiters using agentic AI systems that completely automate entire sections (filtering) of the recruiting workflow: as P_8 states, "*I can almost blindly move those 15%... to talent interview, because I know that they have passed the bot.*" Despite "*blindly*" turning over decisions to the bot, P_8 said she "*I still feel I'm in control... because... I am still the responsible one, so the bot is only gonna do what I tell the bot to do.*"

4.1.2 AI as the invisible architect. Despite participants' descriptions of how little AI and genAI tools influenced their decision process, our interviews demonstrated that genAI systems were invisible architects of the recruitment process. Everything from determining the relevant qualifications for a position, to generating interview questions and rubrics for assessing candidates, to summarizing and highlighting the most relevant pieces of information from an interview to make a final decision, were generated or summarized by an AI system. This fact, plus the common habit of deferring to AI output, means that while the 'decision' about whether to select a candidate is often made by a human, genAI, subtly and incrementally, can significantly shape the hiring process.

For example, P_9 shared how genAI systems were essentially responsible shaping the framework through which recruiters understand what skills the job even requires, and how they should evaluate and reason about applicants:

"Theming, interview themes, ... like talking with a, hiring manager about a job... We'll... transcribe that [conversation] and theme it and get all the information so then we can build the concepts for us, the questions for us, the rubrics based upon the job description, the panel interview questions ... We'll use AI to help, you know,

organize the information and then generate things that are useful for us ... so Gemini will transcribe and theme, when we get it ... we don't have to do too much beyond that" (P₉)

Importantly, P₉ states here that the genAI outputs which frame decision-making don't need "too much" done to them, i.e. little editing or oversight, and can be simply used in the recruiting process. Further, despite saying they "use AI" to create "things that help us assess the candidates in the process" P₉ insists that "In terms of, like, assessing applications, we don't use it as much as... some other organizations do"—downplaying AI's influence.

Beyond AI designing decision-making frameworks, our interviews also showed that the processing of information relevant to assess candidate qualifications was often also left to AI. Several participants shared how they would use AI summaries to process and highlight information gleaned from candidate interviews, and often these summaries got moved straight along to higher ups. For example, P₁₂ stated:

"I use Crisp AI for note-taking in all of my interviews. It does a transcript and an audio recording and has an AI chatbot that you can ask 'summarize all of their answers to the interview questions, and list all the questions they ask me verbatim' ... then I can stick that in there for the hiring manager. (P₁₂)"

Again, P₁₂ "sticks" the AI summaries in the candidates file for the hiring manager with seemingly little oversight. They went on to explain that while the transcripts exist to as a ground truth for what was said during the interview, "it's very uncommon that someone will actually go and review them." P₈ seconded this practice, stating

"I kind of tell the bot what... to summarize, out of this talk. So, for example, I would say, Summarize... the candidate's experience with enterprise-level clients and then the AI bot would understand that, and I can then fill that into the [candidate's] scorecard." (P₈)

Adding to this pattern, P₁₀ states that thanks to the AI summarization, "I don't have to do anything" to further process the interview. Evidenced by these examples, AI and genAI systems have huge influence over recruiting decision-making processes by designing assessment frameworks and processing relevant applicant information—but recruiters do not recognize its influence, despite handing over control to AI.

5 Discussion and Concluding Thoughts

We reflect on mechanisms for reclaiming human agency in recruiting workflows for practitioners, researchers, and policymakers.

Implications for AI policy surrounding decision-making Much law surrounding AI accountability focuses on decision-making outputs: In the U.S., Title VII prevents disparities in employment-related assessment outputs that cannot be justified by "business necessity;" in the EU, the GDPR emphasize the right not to be subject to decisions "based solely on automated processing." [22, 62] However, our findings show that the introduction of genAI systems to high-stakes decision-making pipelines significantly broadens the sphere of influence that AI systems have over these decisions: Generative AI systems create the entire *infrastructure* (see [45, 51])

within which that decision rests: organizational policy documents and rubrics; in addition to *processing* the information that is considered. This creates a tension: Recruiters can claim that a final decision was human authored, even as the decision space itself has been extensively structured by AI. As a result, legal protections around automated *decisions* may fail to capture the ways in which agency are altered by genAI-assisted processes. Whether or not this is a problem is a question, however, our work show that if limiting decision-making based on AI systems is aspirational, our current frameworks may be insufficient.

Implications for research. Several papers have pointed to the fact that genAI systems have fundamentally changed the way that bias testing and mitigation [40] functions in AI systems, and calls into questions whether it is possible [5]. First, our work shows an empirical observation of a more technical point made by many [62, 66]: how does one mitigate bias a textual output with an unclear mapping to the final decision? In the case of recruitment, how does one understand the possibility of AI bias in a generated rubric, or set of interview questions? More generally, as researchers, this points to a growing need to build tools to understand the impacts of intermediate model outputs in a decision-making process, not just those providing a concrete recommendation. This leads to a natural question: how can we help practitioners retain, or achieve, awareness of AI's influence as its role becomes more and more subtle and intertwined in every step of the human thought process towards decision-making?

Implications for organizational practice. Many resources for employees in the workforce are focused on prompt engineering; how to prompt genAI better to produce improved results. This is a tempting response to presumed participant need; if recruiters just figure out how to use genAI, then this will result in time being saved. The opportunity to learn and develop intentionality due to the allure of summarization and saving minutes but losing hours, we suggest that future onboarding should incorporate a form of intentional friction to any genAI use. This may involve pedagogical interventions that force recruiting professionals to articulate their professional logic behind any gut instinct. Such recommendations mirror Bućinca et al.'s [14] recommendations to incorporate disruptive design elements that force a user to explain their own evaluation criteria or professional logic first, prior to seeing automated content. By forcing recruiters to explain why they have a gut feel, this helps to create a reflection-in-action moment; moving from simply checking over genAI work to an independent evaluator.

6 Workshop Contribution

In relation to the 2026 BiAlign Workshop at CHI, we contribute the analysis of a potential loss of agency within a high-stakes decision-making context. The establishment of an invisible architect requires researchers to join together in developing techniques to mitigate the potential harms of such a process. We hope to engage with like-minded researchers and practitioners to consider disruptive design techniques, like those provided by Bućinca et al. [14], as well as possible education-based remediation.

References

- [1] 2026. Origin and history of agency. <https://www.etymonline.com/word/agency>

- [2] Aditya Challapally, Chris Pease, Ramesh Raskar, and Pradyumna Chari. 2025. *The GenAI Divide*. Technical Report. MIT.
- [3] Haozhe An, Christabel Acquaye, Colin Wang, Zongxia Li, and Rachel Rudinger. 2024. Do Large Language Models Discriminate in Hiring Decisions on the Basis of Race, Ethnicity, and Gender? *arXiv preprint arXiv:2406.10486* (2024).
- [4] Jiafu An, Difang Huang, Chen Lin, and Mingzhu Tai. 2025. Measuring gender and racial biases in large language models: Intersectional evidence from automated resume evaluation. *PNAS nexus* 4, 3 (2025), pgaf089.
- [5] Jacy Reese Anthis, Kristian Lum, Michael Ekstrand, Avi Feller, and Chenhao Tan. 2025. The impossibility of fair LLMs. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 105–120.
- [6] Lena Armstrong, Abbey Liu, Stephen MacNeil, and Danaë Metaxa. 2024. The Silicon Ceiling: Auditing GPT's Race and Gender Biases in Hiring. In *Proceedings of the 4th ACM Conference on Equity and Access in Algorithms, Mechanisms, and Optimization* (San Luis Potosi, Mexico) (EAAMO '24). Association for Computing Machinery, New York, NY, USA, Article 2, 18 pages. doi:10.1145/3689904.3694699
- [7] Ashby. 2025. Ashby. <https://www.ashbyhq.com/>
- [8] Alessio Azzutti. 2024. Artificial Intelligence and Machine Learning in Finance: Key Concepts, Applications, and Regulatory Considerations. In *The Emerald Handbook of Fintech: Reshaping Finance*. Emerald Publishing Limited, 315–339.
- [9] Tina Behzad, Siddhartha Devic, Vatsal Sharan, Aleksandra Korolova, and David Kempe. 2025. An External Fairness Evaluation of LinkedIn Talent Search. *arXiv:2511.10752* [cs.CY] <https://arxiv.org/abs/2511.10752>
- [10] Rosanna Bellini, Emily Tseng, Noel Warford, Alaa Daffalla, Tara Matthews, Sunny Consolvo, Jill Palzkill Woelfer, Patrick Gage Kelley, Michelle L Mazurek, Dana Cuomo, Nicola Dell, and Thomas Ristenpart. [n.d.]. SoK: Safer Digital-Safety Research Involving At-Risk Users. (n.d.).
- [11] Virginia Braun and Victoria Clarke. 2021. Thematic analysis: A practical guide. (2021).
- [12] Breezy. 2025. Breezy. <https://breezy.hr/>
- [13] BrightHire. 2025. BrightHire. <https://brighthouse.com/>
- [14] Zana Bućinca, Maja Barbara Malaya, and Krzysztof Z. Gajos. 2021. To Trust or to Think: Cognitive Forcing Functions Can Reduce Overreliance on AI in AI-assisted Decision-making. *Proc. ACM Hum.-Comput. Interact.* 5, CSCW1, Article 188 (April 2021), 21 pages. doi:10.1145/3449287
- [15] Bullhorn. 2025. Bullhorn. <https://www.bullhorn.com/>
- [16] Shiye Cao and Chien-Ming Huang. 2022. Understanding User Reliance on AI in Assisted Decision-Making. *Proc. ACM Hum.-Comput. Interact.* 6, CSCW2, Article 471 (Nov. 2022), 23 pages. doi:10.1145/3555572
- [17] Myra Cheng, Sunny Yu, Cino Lee, Pranav Khadpe, Lujain Ibrahim, and Dan Jurafsky. 2025. Social Sycophancy: A Broader Understanding of LLM Sycophancy. doi:10.48550/arXiv.2505.13995 arXiv:2505.13995 [cs].
- [18] Avishek Choudhury and Zaira Chaudhry. 2024. Large language models and user trust: consequence of self-referential learning loop and the deskilling of health care professionals. *Journal of Medical Internet Research* 26 (2024), e67674.
- [19] Covey. 2024. Covey. <https://covey.framer.website/>
- [20] Maria De-Arteaga, Riccardo Fogliato, and Alexandra Chouldechova. 2020. A Case for Humans-in-the-Loop: Decisions in the Presence of Erroneous Algorithmic Scores. In *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems* (Honolulu, HI, USA) (CHI '20). Association for Computing Machinery, New York, NY, USA, 1–12. doi:10.1145/3313831.3376638
- [21] Eightfold. 2025. AI talent acquisition & recruiting platform. <https://eightfold.ai/>
- [22] European Parliament and Council of the European Union. 2016. General Data Protection Regulation. *Official Journal of the European Union* L 119 (2016), 1–88. Article 22.
- [23] Gem. 2025. Gem. <https://www.gem.com/>
- [24] Kate Glazko, Yusuf Mohammed, Ben Kosa, Venkatesh Potluri, and Jennifer Mankoff. 2024. Identifying and Improving Disability Bias in GPT-Based Resume Screening. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency* (Rio de Janeiro, Brazil) (FAccT '24). Association for Computing Machinery, New York, NY, USA, 687–700. doi:10.1145/3630106.3658933
- [25] ATLAS.ti Scientific Software Development GmbH. 2025. ATLAS.ti Mac. <https://atlasti.com>
- [26] Ben Green and Yiling Chen. 2019. Disparate interactions: An algorithm-in-the-loop analysis of fairness in risk assessments. In *Proceedings of the conference on fairness, accountability, and transparency*. 90–99.
- [27] Ben Green and Yiling Chen. 2019. The Principles and Limits of Algorithm-in-the-Loop Decision Making. *Proc. ACM Hum.-Comput. Interact.* 3, CSCW, Article 50 (Nov. 2019), 24 pages. doi:10.1145/3359152
- [28] Ben Green and Yiling Chen. 2021. Algorithmic Risk Assessments Can Alter Human Decision-Making Processes in High-Stakes Government Contexts. *Proc. ACM Hum.-Comput. Interact.* 5, CSCW2, Article 418 (Oct. 2021), 33 pages. doi:10.1145/3479562
- [29] Greenhouse. 2025. Greenhouse. <https://www.greenhouse.com>
- [30] Hirevue. 2025. HireVue. <https://www.hirevue.com/>
- [31] Lujain Ibrahim, Katherine M Collins, Sunnie SY Kim, Anka Reuel, Max Lam-parth, Kevin Feng, Lama Ahmad, Prajna Soni, Alia El Kattan, Merlin Stein, et al. 2025. Measuring and mitigating overreliance is necessary for building human-compatible AI. *arXiv preprint arXiv:2509.08010* (2025).
- [32] iCIMS. 2025. iCIMS. <https://www.icims.com/>
- [33] Indeed. 2025. Indeed. <https://www.indeed.com/hire>
- [34] Juicebox. 2025. Juicebox (PeopleGPT) - The Leading AI Recruiting Platform. <https://juicebox.ai/>
- [35] Klaus Krippendorff. 2011. Computing Krippendorff's alpha-reliability. (2011).
- [36] Olya Kudina and Bas de Boer. 2025. Large language models, politics, and the functionalization of language. *AI and Ethics* 5, 3 (2025), 2367–2379.
- [37] Alisa Küper and Nicole Krämer. 2025. Psychological traits and appropriate reliance: Factors shaping trust in AI. *International Journal of Human-Computer Interaction* 41, 7 (2025), 4115–4131.
- [38] Vivian Lai, Chacha Chen, Q Vera Liao, Alison Smith-Renner, and Chenhao Tan. 2021. Towards a science of human-ai decision making: a survey of empirical studies. *arXiv preprint arXiv:2112.11471* (2021).
- [39] Lan Li, Tina Lassiter, Joohee Oh, and Min Kyung Lee. 2021. Algorithmic Hiring in Practice: Recruiter and HR Professional's Perspectives on AI Use in Hiring. In *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society* (Virtual Event, USA) (AI/ES '21). Association for Computing Machinery, New York, NY, USA, 166–176. doi:10.1145/3461702.3462531
- [40] Yingji Li, Mengnan Du, Rui Song, Xin Wang, and Ying Wang. 2023. A survey on fairness in large language models. *arXiv preprint arXiv:2308.10149* (2023).
- [41] LinkedIn. 2025. *The Future of Recruiting 2025*. Technical Report. <https://business.linkedin.com/talent-solutions/resources/future-of-recruiting>
- [42] Takuya Maeda and Anabel Quan-Haase. 2024. When Human-AI Interactions Become Parasocial: Agency and Anthropomorphism in Affective Design. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT '24)*. Association for Computing Machinery, New York, NY, USA, 1068–1077. doi:10.1145/3630106.3658956
- [43] Lars Malmqvist. 2025. Sycophancy in Large Language Models: Causes and Mitigations. In *Intelligent Computing*, Kohei Arai (Ed.). Springer Nature Switzerland, Cham, 61–74. doi:10.1007/978-3-031-92611-2_5
- [44] Nestor Maslej, Loredana Fattorini, Raymond Perrault, Yolanda Gil, Vanessa Parli, Njenga Kariuki, Emily Capstick, Anka Reuel, Erik Brynjolfsson, John Etchemendy, et al. 2025. Artificial intelligence index report 2025. *arXiv preprint arXiv:2504.07139* (2025).
- [45] Stephanie Mertens, Mario Herberz, Ulf J. J. Hahnel, and Tobias Brosch. 2022. The effectiveness of nudging: A meta-analysis of choice architecture interventions across behavioral domains. *Proceedings of the National Academy of Sciences* 119, 1 (Jan. 2022), e2107346118. doi:10.1073/pnas.2107346118 Publisher: Proceedings of the National Academy of Sciences.
- [46] Metaview. 2025. Metaview. <https://www.metaview.ai/>
- [47] Huy Nghiem, John Prindle, Jieyu Zhao, and Hal Daumé Iii. 2024. "You Gotta be a Doctor, Lin": An Investigation of Name-Based Bias of Large Language Models in Employment Recommendations. *arXiv preprint arXiv:2406.12232* (2024).
- [48] Phil Strazzulla. 2025. 22 Best Applicant Tracking Systems (ATS): Full Comparison 2025. <https://www.selectsoftwarereviews.com/buyer-guide/applicant-tracking-systems>
- [49] Manish Raghavan, Solon Barocas, Jon Kleinberg, and Karen Levy. 2020. Mitigating bias in algorithmic hiring: evaluating claims and practices. In *Proceedings of the 2020 Conference on Fairness, Accountability, and Transparency (Barcelona, Spain) (FAT* '20)*. Association for Computing Machinery, New York, NY, USA, 469–481. doi:10.1145/3351095.3372828
- [50] LinkedIn Recruiter. 2025. LinkedIn Recruiter. <https://business.linkedin.com/talent-solutions/recruiter>
- [51] Devansh Saxena and Shion Guha. 2024. Algorithmic Harms in Child Welfare: Uncertainties in Practice, Organization, and Street-level Decision-making. *ACM Journal on Responsible Computing* (March 2024). doi:10.1145/3616473 Publisher: ACM/PUB27New York, NY.
- [52] Ben Shneiderman, Catherine Plaisant, Maxine Cohen, Steven Jacobs, Niklas Elmqvist, and Nicholas Diakopoulos. 2016. *Designing the User Interface: Strategies for Effective Human-Computer Interaction* (6th ed.). Pearson.
- [53] S Shyam Sundar. 2020. Rise of Machine Agency: A Framework for Studying the Psychology of Human-AI Interaction (HAI). *Journal of Computer-Mediated Communication* 25, 1 (Jan. 2020), 74–88. doi:10.1093/jcmc/zmz026 _eprint: <https://academic.oup.com/jcmc/article-pdf/25/1/74/32961171/zmz026.pdf>.
- [54] Helena Vasconcelos, Matthew Jörke, Madeleine Grunde-McLaughlin, Tobias Gerstenberg, Michael S Bernstein, and Ranjay Krishna. 2023. Explanations can reduce overreliance on ai systems during decision-making. *Proceedings of the ACM on Human-Computer Interaction* 7, CSCW1 (2023), 1–38.
- [55] Bryan Wilder, Eric Horvitz, and Ece Kamar. 2020. Learning to complement humans. *arXiv preprint arXiv:2005.00582* (2020).
- [56] Willo. 2025. Willo. <https://www.willo.video>
- [57] Christo Wilson, Avijit Ghosh, Shan Jiang, Alan Mislove, Lewis Baker, Janelle Szary, Kelly Trindel, and Frida Polli. 2021. Building and auditing fair algorithms: A case study in candidate screening. In *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*. 666–677.

- [58] Kyra Wilson and Aylin Caliskan. 2024. Gender, Race, and Intersectional Bias in Resume Screening via Language Model Retrieval. *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society* 7, 1 (Oct. 2024), 1578–1590. doi:10.1609/ai.es.v7i1.31748
- [59] Kyra Wilson, Mattea Sim, Anna-Maria Gueorguieva, and Aylin Caliskan. 2025. No Thoughts Just AI: Biased LLM Hiring Recommendations Alter Human Decision Making and Limit Human Autonomy. arXiv:2509.04404 [cs.CY] <https://arxiv.org/abs/2509.04404>
- [60] Workable. 2025. Workable. <https://www.workable.com/>
- [61] Workday. 2025. Workday Talent Management. <https://www.workday.com/en-us/products/talent-management/overview.html>
- [62] Alice Xiang. 2024. Fairness & privacy in an age of generative ai. *Science and Technology Law Review* 25, 2 (2024).
- [63] John R Zech, Marcus A Badgeley, Manway Liu, Anthony B Costa, Joseph J Titano, and Eric Karl Oermann. 2018. Variable generalization performance of a deep learning model to detect pneumonia in chest radiographs: a cross-sectional study. *PLoS medicine* 15, 11 (2018), e1002683.
- [64] Chunpeng Zhai, Santoso Wibowo, and Lily D Li. 2024. The effects of over-reliance on AI dialogue systems on students’ cognitive abilities: a systematic review. *Smart Learning Environments* 11, 1 (2024), 28.
- [65] Shuning Zhang, Hui Wang, and Xin Yi. 2025. Exploring Collaboration Patterns and Strategies in Human-AI Co-creation through the Lens of Agency: A Scoping Review of the Top-tier HCI Literature. *Proc. ACM Hum.-Comput. Interact.* 9, 7, Article CSCW413 (Oct. 2025), 43 pages. doi:10.1145/3757594
- [66] Thomas Zollo, Nikita Rajaneesh, Richard Zemel, Talia Gillis, and Emily Black. 2025. Towards Effective Discrimination Testing for Generative AI. In *Proceedings of the 2025 ACM Conference on Fairness, Accountability, and Transparency*. 1028–1047.

A Participant Table

See Table 1 for de-identified demographic information for our participants.

B Taxonomy of AI-Assisted Recruitment Software

To further investigate the unique tools mentioned by our interviews, we first confirmed that such software were regularly used by other industry professionals (i.e., non-bespoke) with popular aggregators [48]. Then, we identified AI features across each system, then categorized them according to two qualities: what stage of the recruitment pipeline the tool is intended to be used for (*recruitment phase*), and, what type of action being performed (*functionality*). Our results can be shown in Table 2.

In parallel to our interview analyses, we found that there exist AI, and specifically genAI, applications throughout the first four stages of the recruitment process (2), excluding the extension of the final offer. The most common application of AI is ‘Database filtering,’ which facilitates sourcing by searching across candidate databases, though this functionality existed prior to 2022 and the introduction of ChatGPT. ‘Job Description Generation,’ which requires genAI technology, is the second most prevalent application of AI. The stage with the most genAI-specific applications is ‘Interviewing,’ where genAI has introduced new functionalities, such as summarizing interviews and generating interview questions. The AI functionalities that participants did not explicitly mention were used include ‘AI-led screening,’ ‘Candidate filtering,’ ‘Interview scoring,’ ‘Interview feedback,’ and ‘AI-led interview.’

C Interview Protocol

We provide our complete semi-structured interview protocol below. Due to the nature of semi-structured interviews, we asked follow-up questions and omitted other questions in the protocol as needed.

Grounding Questions

Q1 Tell me about your work and your organization.

- (1) How long have you been in this role?

Q2 Can you tell me about how you are involved with company hiring processes, and a bit about what that process looks like (think of a recent hire)?

- (1) Are there steps you normally go through? Frequency? Number of people involved?
- (2) Is there a rubric you use to evaluate people? What factors do you often consider?

General Use of GenAI in Hiring

Q1 Have you used an LLM to support you with your work related to hiring? If so, how so and where in the process do you use it?

- (1) ChatGPT, Claude are popular tools, what model do you use?
- (2) Is this a general model or personalised to your company?
- (3) Is this a personal or business account?
- (4) What motivated you to use it?

Q2 What do you ask the LLM to do <in the hiring process stated previously>?

- (1) Are there specific tasks you assign to this?
- (2) What information do you give it?
- (3) What do you do with the output?
- (a) Is this useful to you? If yes/no, why?

Honing in on One Application of GenAI

Q1 What are the inputs/outputs at this part of the process?

- (1) How do you use the outputs to reach the next stage of the decision process?

Q2 Has your use of AI matched with your expectations?

Q3 You mention <AI-assisted task> - have you ever looked at the <AI output, e.g. summary of an interview> and felt that it changed your impression of <the input, e.g. a candidate/the position>?

- (1) Has it brought things to your attention you wouldn’t have noticed on your own?

Q4 Do you find yourself deferring to the <AI output>?

- (1) Do you think you do so just because it’s convenient?
- (2) Do you think these tools are improving the selection process?
- (3) What factors make you trust or question its judgment even when your professional instinct disagrees?

Q5 Do you ever check for errors on the outputs of the LLM? What do you look for?

Perception of Agency over Choices

Q1 Is there a reason you haven’t used LLMs in other parts of your hiring process?

- (1) Is there something you would never ask an LLM to do for you?
- (2) Do you have any concerns about how it is currently used or could be used in your company?

Q2 Do you feel like you understand why it produces the candidates, scores, summaries etc. that it does?

- (1) Has there been anything that challenges this perspective?

Q3 How much do you feel like you are still in control of the final decision, versus the system guiding or constraining your choices?

Q4 Do you feel like you've lost any skills as a result of relying on AI?

- (1) If the AI suddenly stopped working tomorrow, would you know how to tackle something?

Q5 How do you explain your hiring choices to others (e.g., managers) when the AI system has played a role in shaping those choices?

Q6 <If participant uses genAI in applicant evaluation> Can you give me an example of a prompt you've used to evaluate applicants with an LLM? This prompt can be generic.

D Codebook

See Table 3 for the codebook resulting from our process of thematic analysis.

Received 25 March 2026

<i>ID</i>	<i>Age</i>	<i>Gender</i>	<i>Most Recent Job Title</i>	<i>Exp.</i>	<i>Tools Discussed</i>
P_1	–	–	Chief Executive Officer	–	ChatGPT, n8n
P_2	35-44	Man	Sr. Recruitment Specialist	17	Internal system
P_3	35-44	Man	Director of Recruitment	15	BrightHire, Ashby, Gemini, Gem
P_4	–	–	Talent Acquisition Specialist	–	Indeed, ChatGPT
P_5	–	–	Director of Recruitment	–	Metaview, Internal tool, Ashby, Granola, ChatGPT, CoderPad
P_6	–	–	Chief Operating Officer	–	LinkedIn, ChatGPT
P_7	35-44	Woman	Director of Recruitment	22	LinkedIn, ChatGPT
P_8	35-44	Woman	Sr. Recruitment Specialist	10	Undisclosed AI agent service, Greenhouse, Mettl
P_9	45-54	Man	Director of Recruitment	27	Workday, Gemini
P_{10}	35-44	Woman	Talent Acquisition Specialist	13	Paraform, LinkedIn, Ashby, Juicebox, Metaview, ChatGPT
P_{11}	25-34	Woman	Sr. Human Resources Manager	8	ChatGPT, LinkedIn, Zoom
P_{12}	25-34	Man	Talent Acquisition Specialist	5	Willo, Greenhouse, Crisp
P_{13}	–	–	Talent Acquisition Specialist	–	ChatGPT, Gemini
P_{14}	35-44	Woman	Talent Acquisition Specialist	14	LinkedIn, Internal LLM, ChatGPT, Google Meet
P_{15}	35-44	Woman	Sr. Recruitment Specialist	15	ChatGPT, Greenhouse, LinkedIn
P_{16}	35-44	Woman	Talent Acquisition Specialist	6	ChatGPT, LinkedIn
P_{17}	25-34	Man	Talent Acquisition Specialist	4	ChatGPT, LinkedIn
P_{18}	–	–	Talent Acquisition Specialist	–	ChatGPT, LinkedIn
P_{19}	35-44	Woman	Talent Acquisition Specialist	11	ChatGPT, LinkedIn, Workday
P_{20}	35-44	Woman	Sr. Recruitment Specialist	15	HireVue, iCIMS, Copilot
P_{21}	18-24	Woman	Talent Acquisition Specialist	1	Copilot, Internal system, LinkedIn
P_{22}	18-24	Woman	Talent Acquisition Specialist	2	Bullhorn, ChatGPT, LinkedIn

Table 1: Interview participant demographic information. “Most Recent Job Title” refers to a general class of job titles to retain anonymization. “Tools Discussed” only includes tools that were actively used by the participant, not merely mentioned in their interview.

	Job description generation	Database filtering	Candidate summary	AI-led screening	Candidate filtering	Qualification grouping	Candidate fraud detection	Match ranking/score	Interview summary	Interview question generation	Interview transcription	Interview scoring	Interview feedback	AI-led interview
Applicant Tracking Systems														
Workable [†] [60]	•	•	•				•		•					
Ashby [†] [7]	•	•		•		•				•				
Gem ^{†*} [23]		•		•			•	•						
Breezy [†] [12]						•		•						
Bullhorn [†] [15]	•		•	•			•	•	•					
iCIMS ^{†*} [32]	•	•					•		•					
Greenhouse [†] [29]	•			•				•	•					
Workday [†] [61]	•	•		•	•									
Supplemental Tooling														
Indeed [†] [33]	•	•	•				•							
LinkedIn Recruiter ^{†*} [50]		•	•		•									
Metaview ^{†*} [46]	•	•						•		•				
HireVue [†] [30]								•	•	•	•			
BrightHire [†] [13]	•			•				•	•	•		•		
Willo [†] [56]				•				•	•	•	•			
Juicebox ^{†*} [34]		•	•											
Agentic Software														
Eightfold [21]	•	•		•			•	•		•	•			•
Covey [19]		•					•	•						

Table 2: Software is organized the AI functionalities it supports and ordered by recruitment stage (Section 2). Applicant tracking systems refer to platforms that support the full-cycle recruitment pipeline, while supplemental tooling offers tooling that focus on a single pipeline phase. Agentic software or agents refer to the distinct functionality of autonomous action that can operate with minimal human intervention (identified by *). For clarity, we denote software that participants reported to actively use with [†]

Company-focused	Candidate-focused
<i>Business Culture</i> : What is the company's culture like? <i>Business Size</i> : How large is the company <i>Pressure to Use AI in Business</i> : Participant reporting pressure to use AI from their employer <i>Stage of AI Adoption</i> : How much AI the employer is incorporating company-wide <i>Prohibitive Cost of AI</i> : Statement that AI tools are out of the price range for a participant's company	<i>Candidate Use of AI</i> : How candidates use AI <i>Candidate Effort stands out</i> : The manual effort of a candidate is obvious when reviewing their application <i>Desired Candidate Skillsets</i> : Skills the participant looks for <i>Applicant-recruiter Arms Race</i> : Participant's use of AI to combat candidate's use <i>Fake Candidates</i> : Discussion of reportedly "fake" candidates <i>High Volume of Applications</i> : Capturing the great demand on recruiters for reviewing applications
Agency-focused	Impacts/Future
<i>Auto/AI-Resistant Work</i> : Assertion that a specific task cannot be performed by AI <i>Deference to AI</i> : Instance where the participant trusts AI <i>Employer AI Red Line</i> : Determination of when AI should not be used in the recruitment process <i>Human Agency Maintained</i> : Belief that the participant maintains decision-making agency <i>Lack of Trust in AI/Skepticism</i> : Skepticism toward the use of AI in general or a specific task <i>Shared Decision-making</i> : Belief that AI has some impact on the decision-making process <i>Perceived Human Manual</i> : Claim that participant is manually performing a certain task	<i>AI Future Potential</i> : Looking forward to AI's future impact <i>Legal and Privacy Risk</i> : Reported fear of risk or performing actions that constitute risk <i>Loss of Skill/Deskilling</i> : Observed behavior that reflects deskilling <i>Overpromising of AI</i> : Belief that AI marketing has overpromised on the utility of AI <i>AI Used Differently by Generations</i> : Pattern that different age groups use AI in disparate manners <i>Two-tiered AI use</i> : Different usage of AI depending on qualities of the role or candidate
Participant/Role Information	Recruitment Stage
<i>Job Level Type</i> : The level of position that the participant hires <i>Participant Background</i> : Information about participant's role and experience <i>Self-Assessment of AI Proficiency</i> : A participant's confidence in their use of AI <i>Time in Role</i> : How long has the participant been in this role	<i>Assessment</i> : Process of forming a judgment on a candidate <i>Conceptualizing/Themeing</i> : Generation of a job description <i>Final Decision-making</i> : Selection of candidate(s) to send an offer <i>Interviewing</i> : Conversations with candidates <i>Logistics</i> : Planning interviews, sending communication, etc. <i>Sourcing</i> : Finding candidates to apply to a job opening <i>Screening</i> : Comparing a candidate's skills and qualifications to the job requirements
Task Type	Performance of AI Systems
<i>Agentic AI</i> : Using AI agents to complete a task <i>AI Content Generation</i> : Using genAI to produce summaries, etc. <i>Analogue-Technical Processes</i> : A process combining both technical and manual methods <i>Grouping</i> : Classifying candidates in a fixed number of groups <i>Keyword Searching</i> : Looking for specific keywords in an input <i>Manual Work</i> : Human-only completion of a recruitment task <i>Transcription</i> : Translation from a recording to text <i>Summarizing</i> : Reduction of an input into a shortened form <i>Ranking</i> : Comparing candidates to one another individually <i>Research/Information Gathering</i> : Process of learning about a certain domain <i>System Indicators of Job Compatibility</i> : How is job compatibility reflected to the participant	<i>AI Evaluates Candidates Badly</i> : Inconsistent or incorrect evaluation of candidates <i>AI is Biased</i> : Relaying of concerns about biased outputs from AI <i>AI Lacks Nuance</i> : Discusses that AI systems may be brittle or do not understand humans <i>AI Makes More Work</i> : Claim that the participant's use of AI increases their workload <i>AI Reveals Overlooked Details</i> : AI provides the participant with a more detailed interpretation of the input <i>AI Saves Time</i> : The use of AI reduces time spent on a task <i>AI Tool Performance</i> : Reflects how well the AI tool is seen to perform in general <i>AI Use is Essential for Efficiency</i> : Belief that one's use of AI is essential for the efficient completion of tasks

Table 3: The codebook derived from our thematic analysis, grouped by the codes' focus.